

B. TECH 5TH SEMESTER EXAMINATION, DECEMBER 2021

Database Management Systems [CS 3102]

Name: Abhiroop Mukherjee

Examination Roll No.: 510519109

G-Suite ID : 510519109.abhirup@students.iests.ac.in

Q4) a) Physical Data Independence.

✓, 2, 3, 4, 5

→ The capacity of a DBMS to allow change of internal schema without affecting conceptual or ~~external~~ external schema is called the Physical Data Independence.

Eg: consider we change HDD of a computer running a DBMS to a bigger size with all previous data, this should not affect any schemas or views in ~~DB~~ DBMS.

b) Cardinality

→ cardinality ~~with~~ with respect to DBMS refers to the uniqueness of data value contained in the column

High Cardinality → most of the data in a column are unique

low cardinality → there are a lot of duplicate data in a column

c) Primary Key

→ before going into primary key, let's define what a key is

→ If $R(A_1, A_2, \dots, A_n)$ is a relation with ~~pr~~ functional dependencies F & X is a subset of R , we say X is a key if

i) $X \rightarrow A_1, A_2, \dots, A_n$ is implied in F

ii) R is not functionally dependant on any proper subset of X

→ There can exist multiple keys for a relation. The set of those keys are called candidate key

→ The key chosen by Database designed from the set of candidate keys ~~are~~ is called the primary key.

d) 3NF

A relation scheme R is in 3NF, if whenever $X \rightarrow A$ holds in R & A is not in X , then either X is superkey of R or A is a prime attribute.

ie 3NF is 2NF + constraint (no \rightarrow nonprime \rightarrow nonprime)

c) Lock compatibility matrix

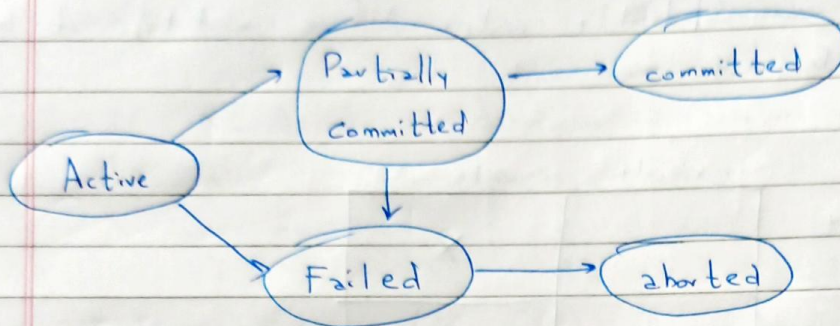
→ Lock compatibility matrix defines what type of locks can be made over each other.

	read only	write only
read-only		X
write only	X	X

→ This matrix means that ~~then a transaction~~ if a data is read-locked by a Transaction T_i , some other transaction T_j can also put a read lock on it (but not write lock)

→ but if data is write locked by T_i , ~~by~~ T_j can't put read nor write lock on it.

Q3) 2) State Transition Diagram of a Transaction

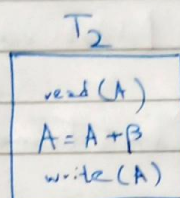
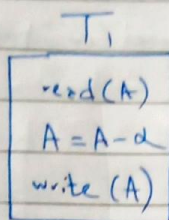



- Partially committed state of a transaction means that the execution of transaction is done ~~and if all goes~~ ~~DBMS~~
- If all DBMS checks is passed by the transaction, then the transaction goes into committed state. Changes made by the transaction is now permanent
- If for some reason the DBMS checks fail, the ~~trans~~ transaction goes into Failed State, all its changes are rolled back via recovery manager & then the transaction is aborted
- User of DBMS can not redo the transaction after necessary changes so that update to DBMS ~~can be~~ happens.

Q3)b)

Lost Update Problem

consider two transactions



[$d, \beta \rightarrow$ some constant]

→ when T_1 & T_2 happens parallelly, this could happen
[assume initially $A = 1000$]

T_1	T_2	DB MS data.
read(A) [$A = 1000$]	read(A) [$A = 1000$]	$A = 1000$
$[A = 1000 - d]$ $A = A - d$ ↓	$A = A + \beta$ [$A = 1000 + \beta$]	
write(A)	write(A)	$A = 1000 - d$ $A = 1000 + \beta$
		↓ end result $\Rightarrow A = 1000 + \beta$

→ here we see that ~~that~~ the observed result is different from what was expected, update of T_1 is lost

→ Here ~~consist~~ consistence of a Database is compromised, loss of ~~the~~ update immediately puts a database to an inconsistent state. [C of ACID]

Q3) c) Wait-Die Scheme

→ Wait-Die scheme is a method to solve deadlock which can be described as follows

[didn't use "avoidance" nor "prevention" here for Q3d]

i) If a transaction

→ This scheme takes place when a transaction requests to lock a resource which is already held by some other transaction.

→ If that case arrives, we have two possibility

i) If T_i , which is requesting a conflicting lock is older than T_j , i.e. $TS(T_i) < TS(T_j)$.

→ Then T_i waits for T_j to ~~release~~ release ^{it's} ~~this~~ lock then locks the resource for itself.

ii) If $TS(T_i) > TS(T_j)$ [T_i younger than T_j], then T_i is killed and restarted again sometime with same ~~time~~ time stamp as before.

TLDR: older transaction ~~wait~~ waits, but newer transaction is killed.

d) → Wait-Die Scheme is a deadlock prevention scheme.

→ This is due to the fact that this scheme does it work after the transaction is issued & it either kills / wait an already running Transaction.

→ Deadlock prevention, on the other hand actively looks into transaction instruction before even executing to ~~see~~ find if any deadlock could occur or not

Eg: Wait-for graph creates a "dependency graph" and actively tries to avoid loops in dependency graph to avoid ~~loops~~ deadlocks.

Q5) 2) Multivalued Dependency

→ A to

(i) → for a dependency $A \twoheadrightarrow B$, if for a single value $a \in A$, if there exist multiple values of $b \in B$, then the functional dependency $A \rightarrow B$ is said to be a multivalued dependency.

(ii) → For a multivalued dependency to exist, we must need at ~~least~~ least columns to exist in a relation.

Why? → one column relation can't have ^{functional} dependency

→ two column relation can have functional dependency but can't have multivalued dependency

→ If we try to add multivalued dependency, properties of functional dependency will get violated.

Eg consider relation

A	B
a	1
b	2
c	3
d	4

→ this has an fd $A \rightarrow B$, meaning that if ~~the~~

$$t_i(A) = t_j(A) \Rightarrow \text{implies } t_i(B) = t_j(B)$$

→ if we try to ^{make} ~~insert~~ multivalued dependency by adding (2,4), it will violate the definition of functional definition

(iii) for a three attribute relation $R(A, B, C)$, if there is a multivalued dependency between $A \& B$ and $A \& C$, then $B \& C$ should be independent of each other.

→ A relation R having a multivalued dependency can be in BCNF form, but it cannot be in 4NF form

Q5) b) Undo

→ undo'ing a Transaction means that we restore all the data items edited by the transaction to its previous value before ~~trans~~ the start of execution

Eg consider a transaction

T	
read(A)	
$A = A - 50$	
write(A)	
read(B)	
$B = B - 100$	
write(B)	

and consider ~~it~~ before execution of T, $A = 1000$, $B = 500$

So	T	Data.
	read(A) [$A = 1000$]	$A = 1000$, $B = 500$
	$A = A - 50$ [$A = 950$]	
	write(A)	$A = 950$, $B = 500$
	read(B) [$B = 500$]	
	$B = B - 100$ [$B = 400$]	
	write(B)	$A = 950$, $B = 400$

→ now after execution of T, if we undo T, we will have to revert A to 1000 & B to 500.

Redo

→ redo'ing a transaction means setting the value of all data item updated by transaction to the new values.

Eg: consider the same transaction and same ~~data~~ data A & B

→ redo'ing T means we set ~~A & B~~ A = 950 ~~&~~ & B = 400

→ Redo operation is idempotent, ~~this means that if we redo~~

Eg: if we redo T multiple times, it will ^{have} ~~be~~ same effect as redoing T once, i.e. A = 950 & B = 400 will be ~~the~~ constant even after multiple redos.

→ Redo & Undo operation are used in ^{Log based} Crash Recovery to maintain the ACID [~~consistent~~ ~~availability~~] [consistent & durable] property of a Database.