

SIGNALS : HIGHER LEVEL ECF

SIGNAL is a mechanism, a higher level form of ECF, that allows processes and kernel to send a small message that an event of some type has occurred in the system.

- The kernel handles lower level h/w exceptions and are not normally visible to the user process. Signals providing a mechanism to expose those exceptions to the user processes.
- For example; kernel sends the SIGFPE (or floating point error) to the process when it is trying to divide anything by zero.
- When a process is running in foreground and you type 'CNTRL+c' – kernel sends SIGINT signal to each process in the process group.
- Some signal correspond to higher level s/w event; like a SIGKILL may be sent by a user process to kill another process
- A process can send a signal to itself

LINUX SIGNALS

There are about 15 SIGNALS in Unix System V; 30 in LINUX of which the first 15 are same as System V.

No	Name	Default Action	Corresponding event
1	SIGHUP	Terminate	Terminal line hangup
2	SIGINT	Terminate	Interrupt from keyboard
4	SIGILL	Terminate	Illegal instruction
6	SIGABRT	Terminate & dump core	Abort signal
8	SIGFPE	Terminate & dump core	Floating point error
9	SIGKILL	Terminate	Kill program
10	SIGUSR1	Terminate	User defined signal 1
17	SIGCHLD	Ignore	A Child stopped/terminated
18	SIGCONT	Ignore	Continue process if stopped
19	SIGSTOP	Stop untill next SIGCONT ¹	Stop signal not from terminal
23	SIGURG	Ignore	Urgent condition on socket
30	SIGPWR	Terminate	Power failure

¹This signal can neither be caught nor ignored

SIGNALS : Sending and Receiving

The transfer of a signal to a destination process is done through two distinct steps

- **Sending a signal:** The kernel sends a signal to a destination process – this delivery takes place for one of the two reasons;
 - i) Detection of a system event (say, divide by 0) or the termination of the child process
 - ii) A process has invoked *kill* function explicitly requesting the kernel to kill any process.

A number of mechanisms are used to send signals that relies on the notion of the *process group*. The function *pid_t getpgrp(void)* is used to get the process group (all child of a parent belong to the same process group of its parent). The function *int setpgid(pid_t pid, pid_t pgid)* may be used to change the process group of *pid* to *pgid*.

SIGNALS : Sending and Receiving

- **Receiving a signal:** A destination process receives a signal when the kernel forces it to react on the delivery.
 - i) The process can ignore the signal,
 - ii) terminate; or
 - iii) catch it by executing a user level function; i.e., the *signal handler*.

SIGNALS : Pending

A signal sent but yet to be received is known as a *pending signal*.

- At any point of time there can be only one signal pending of a particular type; so, if a signal of type k is pending and another signal of the same type comes – it would simply be discarded.
- A process can selectively blocks some signals; the blocked signal can be sent but delivered only when that signal is unblocked
- A pending signal can be received at most once
- the kernel keeps, for each process, a pending and a blocked signal bit vector
- The kernel sets bit k for a signal of type k being delivered in *pending* and clears the same bit k when that signal is received


Sending Signals

The sending mechanism of a signal relies on the notion of a process group.

- Every process belongs to a single process group – which can be determined by *pid_t* `getpgrp(void)`.
- By default a process belongs to the process group of its parent. However, this can be changed by making a call `setpgid(pid_t pid, pid_t pgid)` which puts the process group of the process *pid* to *pgid*. [Note that if *pid* = 0 then the pid of the current process is used]
- And a call `setpgid(0, 0)` from the process, say, 10000 would put process 10000 in a new process group 10000.
- A command-line `/bin/kill -9 12345` sends SIGKILL to the process 12345. And `/bin/kill -9 -12345` sends the SIGKILL signal to every process in the process-group 12345.

More on process group: Foreground and background processes

- The shell creates separate process group for each job (an abstraction of the processes that are created to evaluate a single command line; e.g. *unix > ls|sort < enter >*). Here two processes are created and connected through UNIX pipe. There can be one foreground job and 0 or more background jobs.
- A process group-id is taken from one of the parent processes for a job.



Example of a Shell with one foreground and two background jobs.

Sending SIGNALS

From the keyboard: Typing 'CNTRL+c' from the keyboard results in sending signal SIGINT to every process in the foreground process group. In the default case it terminates the foreground job.

Signal to kill : Processes send signal to other processes (including themselves) by the kill function *int kill(pid_t pid, int sig)*. Note

- If successful, the call returns 0; for error -1.
- $pid = +ve$ sends *sig* to the process numbered *pid*
- $pid = 0$ sends *sig* to every processes in the process group of the parent process *pid*
- $pid = -ve$ sends *sig* to every processes in the process group $|pid|$

Sending SIGNALS with alarm function: A process can send SIGNAL to itself by *unsigned int alarm(unsigned int secs)*

Receiving SIGNALS

When the kernel switches from a process 'p' from kernel mode to user mode (on completion of a system call or context switch etc.) it

- checks the the set of unblocked pending signals (pending & blocked) for 'p'
- if this set is empty (the usual case) the kernel passes control to the next instruction I_{curr} ; however if the set is not empty
- the kernel forces a signal 'k' (usually with the lowest 'k') to the process
- the process may in turn take appropriate action, if necessary even overriding the default action using the *signal* function.
- The default action for the signals mentioned earlier are available in the table
- Also note that the default action for the signals, SIGSTOP and SIGKILL can not be changed.

```
#include <signal.h>
typedef void (*sighandler_t) (int);
sighandler_t signal(int signum, sighandler_t handler);
```

The signal function can change the action of the signal in a couple of ways

- if this set is empty (the usual case) the kernel passes control to the next instruction I_{curr} ; however if the set is not empty
- the kernel forces a signal 'k' (usually with the lowest 'k') to the process
- the process may in turn take appropriate action, if necessary even overriding the default action using the *signal* function.
- The default action for the signals mentioned earlier are available in the table
- Also note that the default action for the signals, SIGSTOP and SIGKILL can not be changed.

```
#include <signal.h>
typedef void (*sighandler_t) (int);
sighandler_t signal(int signum, sighandler_t handler);
```

Receiving SIGNALS: An example

Typing 'CNTL+c' results in termination of the foreground job

```
#include <stdio.h>
#include <signal.h>
#include <unistd.h>
#include <stdlib.h>

void signal_handler(int sig)
{ printf(" Hello \n");
  exit(0);
}

int main() { if (signal(SIGINT, signal_handler)==SIG_ERR)
  perror("error:signal");
  pause(); // wait for the handler to catch the signal
  return 0;
}
```

User defined signal: SIGUSR1

```
#include <stdio.h>
#include <signal.h>
#include <unistd.h>
#include <stdlib.h>

void signal_handler(int sig)
{
    printf("Handler: Hello\n");
}

int main() {
    if (signal(SIGUSR1, signal_handler)==SIG_ERR)
        perror("error:signal");
    printf("Main\n");
    raise(SIGUSR1); // send a signal to the calling process
    return 0;
}
```

User defined signal: SIGUSR1

```
void sig_handler_parent(int signum){
    printf("Parent : Received a response signal from child \n");
}

void sig_handler_child(int signum){
    printf("Child : Received a signal from parent \n");
    sleep(1); kill(getppid(),SIGUSR1);
}

int main(){
    pid_t pid;
    if((pid=fork())<0){
        printf("Fork Failed\n");exit(1);
    }
    else if(pid==0){ /* Child Process */
        signal(SIGUSR1,sig_handler_child); // Register signal handler
        printf("Child: waiting for signal\n");pause();
    }
    else{ /* Parent Process */
        signal(SIGUSR1,sig_handler_parent); // Register signal handler
        sleep(1); printf("Parent: sending signal to Child\n");
        kill(pid,SIGUSR1); printf("Parent: waiting for response\n"); pause();
    }
    return 0;
}

// include proper .h files
```