

Module 3: Transport Layer (Lecture – 3)

Dr. Nirnay Ghosh


Assistant Professor

Department of Computer Science & Technology

IIST, Shibpur

TCP: Round Trip Time Estimation & Timeout

- TCP uses a **timeout/retransmit mechanism** to **recover** from **lost segments**
- **Implementation challenge**: how much **larger** the **timeout** should be than the connection's **round trip-time (RTT)**?
- **Estimating RTT**
 - **SampleRTT**: amount of **time** elapsed between **when** the segment is sent and the **corresponding acknowledgement** is received
 - It is measured for the segments which have been transmitted once
 - It **fluctuates** from **segment to segment** due to **congestion in the routers** and **varying loads on the end systems**
 - Estimation: take weighted combination of the previous value **of EstimatedRTT** and the new value of **SampleRTT**
 - **Exponential Weighted Moving Average (EWMA)**: weight of a given **SampleRTT** decays exponentially fast as the updates proceeds



RTT Samples and RTT Estimates

- **EstimatedRTT**: smoothens the variations in the **SampleRTT**
- **DevRTT**: Variability of the RTT
 - Estimates of how much **SampleRTT** typically **deviates** from the **EstimatedRTT**
 - EWMA of the difference between **SampleRTT** and **EstimatedRTT**


$$\text{DevRTT} = (1 - \beta) \cdot \text{DevRTT} + \beta \cdot |\text{SampleRTT} - \text{EstimatedRTT}|$$

$$\text{EstimatedRTT} = (1 - \alpha) \cdot \text{EstimatedRTT} + \alpha \cdot \text{SampleRTT}$$

TCP: Setting and Managing the Retransmission Timeout Interval

- TCP's timeout interval: should be **greater than or equal** to *EstimatedRTT*
- Otherwise: unnecessary retransmission would be sent
- Timeout interval: **should not be too large** than *EstimatedRTT*
 - TCP **would not quickly** retransmit the segment leading to **larger data transfer delays**
- Desirable to set the time equal to the *EstimatedRTT* plus some **margin**
- Margin:
 - **Large**: if there is a **lot of fluctuation** in the *SampleRTT* values
 - **Small**: if there is **little fluctuation**
- TCP's method for determining the retransmission timeout interval

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 \cdot \text{DevRTT}$$



RTT Samples and RTT Estimates

- Initial *TimeoutInterval*: **1 second** (recommended)
- Subsequently the value is **doubled** to avoid **premature timeout** occurring for a **segment** that will **soon be acknowledged**
- As soon the segment is received, *EstimatedRTT* is updated
- *TimeoutInterval* is computed using the **above formula**


TCP: Reliable Data Transfer

- Three major events related to data transmission and retransmission in the TCP sender are:


- Data received from the upper-layer application
- Timer timeout
- ACK receipt

- Commonly occurring scenarios related to retransmission

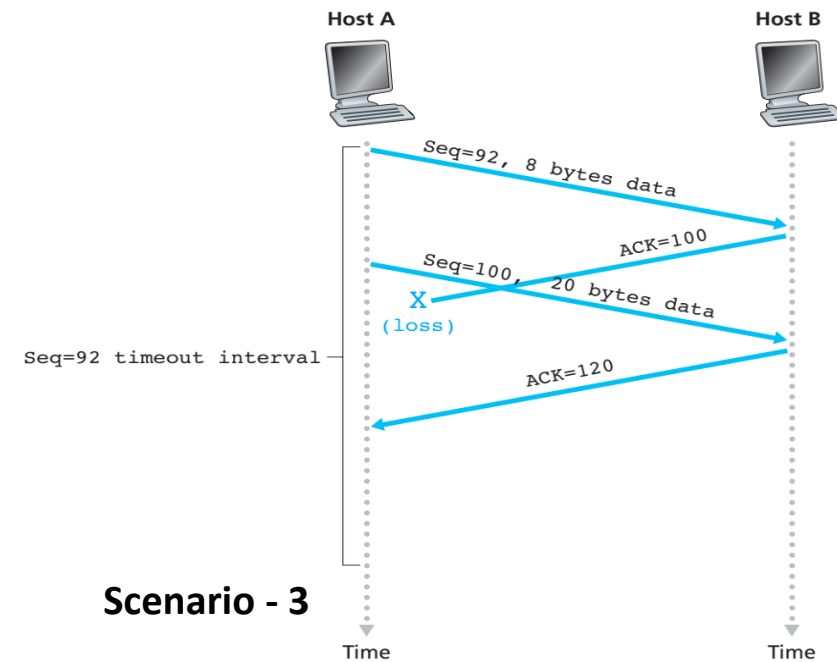
- 1. Retransmission due to lost acknowledgment
- 2. Acknowledgement arrives after the timeout
- 3. Cumulative acknowledgement avoiding retransmission



Scenario - 1



Scenario - 2



Scenario - 3

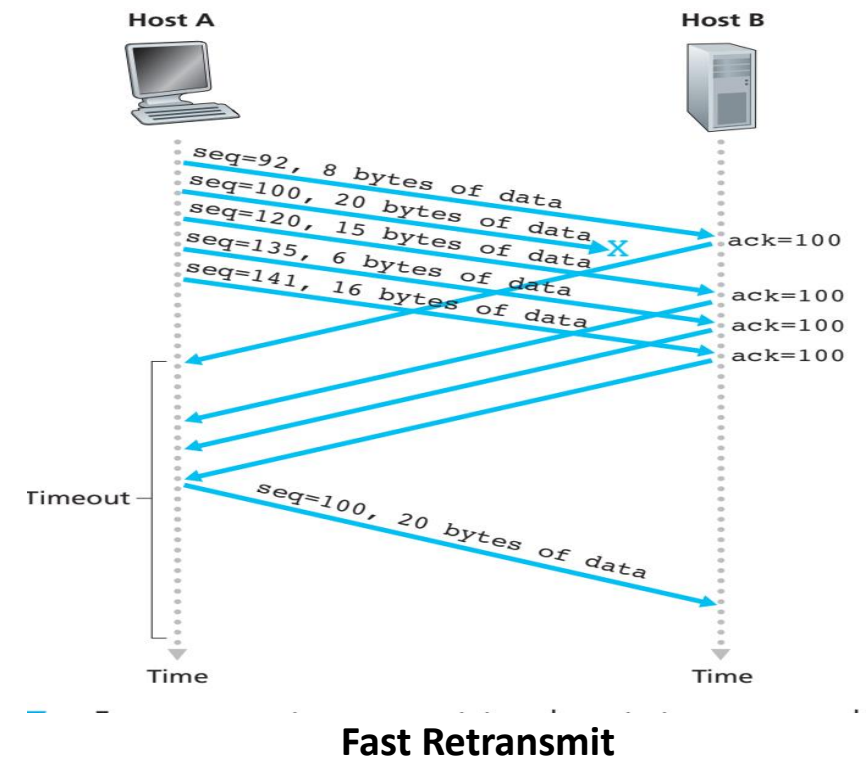
TCP: Reliable Data Transfer – Two approaches

- Doubling the *TimeoutInterval*

- Time out event occurs **first time**: TCP retransmits the **not-yet-acknowledged segment** with **smallest sequence number**
- For **each retransmission**, the *TimeoutInterval* is set to **twice the previous value**
- Timer is **started** under two **events**: data received from application above and ACK received

- *TimeoutInterval* – derived from the most recent values of *EstimatedRTT* and *DevRTT*
- Provides a limited form of **congestion control**
- Congestion: caused due to **persistent retransmissions**
- TCP avoids congestion by **retransmitting after longer and larger intervals**

2/17/2022




- Fast Retransmit

- Longer *TimeoutInterval* forces the sender to delay resending the lost packet
- Increases the end-to-end delay
- **Duplicate ACK**: used by senders to detect packet loss well before the timeout event occurs
- **TCP receiver**: detects a **gap** in the **data stream** – result of **segment loss** or **reordered within the network**
 - Reacknowledges the missing segment multiple times (duplicate ACK) for which the sender has already received an earlier ACK
- **TCP sender**: if receives **three duplicate ACKs** for the same data, performs a **fast retransmission** of the missing segments before the **segment's timer expires**

TCP: Connection Management

- Suppose a **process** running in **one host (client)** wants to **initiate a connection** with **another host (server)**
- **Three-way handshake protocol**: TCP connection establishment involves exchange of three segments:
 - Client-side TCP: sends segment (no application data) with **SYN bit set to 1**; randomly chooses an **initial sequence number (*client_isn*)**; inserts *client_isn* in the **sequence number field** of the TCP segment
 - Server-side TCP: allocates TCP buffers and **variables** to the **connection**; sends a **connection-granted segment** (no application data) with - **SYN bit: 1**, **acknowledgement number field: *client_isn + 1***; **sequence number field: initial sequence number (*server_isn*)**

2/17/2022




Connection Establishment: TCP Three-way Handshake Protocol

- Client-side TCP: allocates **buffers and variables** to the connection; sends **another segment** (may or may not carry application data) with - **SYN bit: 0** (connection already established); **acknowledgement number field: *server_isn + 1***, **sequence number field: *client_isn + 1***

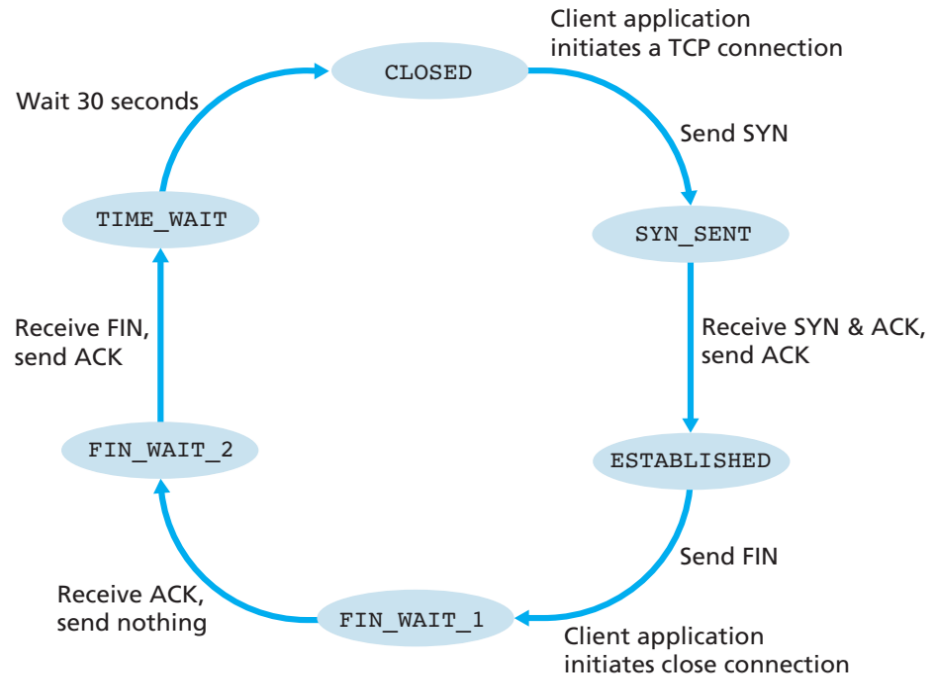
TCP: Connection Management

- Either of the two processes participating in a TCP connection can **end** the connection
- The resources (buffers and variables) in the hosts are **deallocated**.
- Steps involved in closing a TCP connection:
 - **Client TCP** sends a special **TCP segment** to the **server** with **FIN bit** set to 1.
 - The **server TCP**, on receiving this segment, returns an **acknowledgement**.
 - The **server** then sends its own **shutdown segment**, which has the **FIN bit** set to 1.
 - Finally, the client **acknowledges** the server's **shutdown segment**
- At this point, **all the resources** in the **two hosts** are **deallocated**.

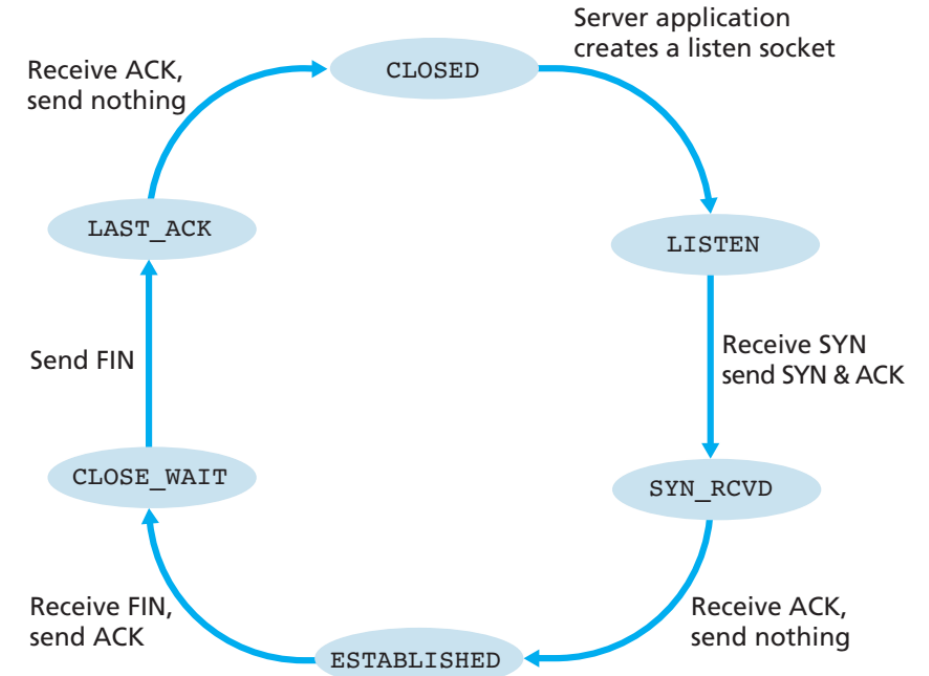


Closing a TCP Connection

TCP States



TCP States visited by a Client-TCP

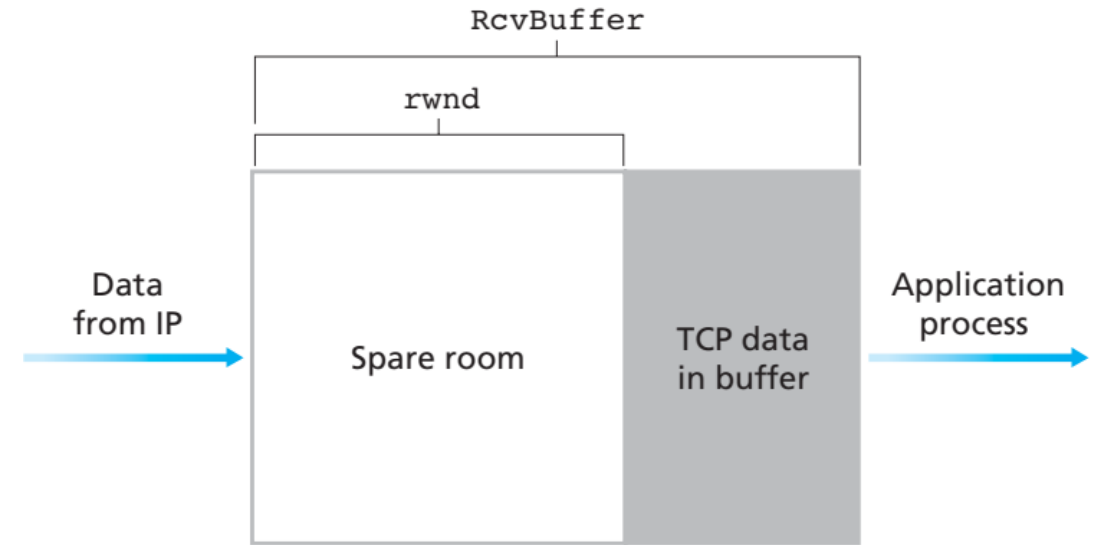


TCP States visited by a Server-TCP

- TCP states visited by a client TCP: **CLOSED, SYN_SENT, ESTABLISHED, FIN_WAIT_1, FIN_WAIT_2, TIME_WAIT**
- TCP states visited by a server TCP: **CLOSED, LISTEN, SYN_RCVD, ESTABLISHED, CLOSE_WAIT, LAST_ACK**
- If host receives a TCP segment whose **port number does not match** with any of the **ongoing sockets in the host**
 - Host sends a special TCP segment with the **RST flag set to 1**
 - Implications: SYN segment reached the target host, but the target host is not running and application with specific port number

TCP: Flow Control

- TCP connection uses **send and receive buffers** to obtain/pass the **data** from/to the **upper-layer sender/receiver application processes**
- The application at the receiver may be busy and read the data after longer duration
- A fast sender can very easily **overflow the connection's receive buffer**
- **Flow control**: eliminates the possibility of the sender overflowing the receiver's buffer
 - Matches the **rate at which the sender is sending** against the **rate at which the receiving application is reading**
- **Sender maintains a variable – receive window (*rwnd*) – initially equal to the receive buffer**
 - The **current value** of *rwnd* is inserted into every **TCP segment** sent by the **receiver**
 - Sender maintains a **distinct receive window** for **each connection**
 - It comes to know of **how much free buffer space** is available at the receiver

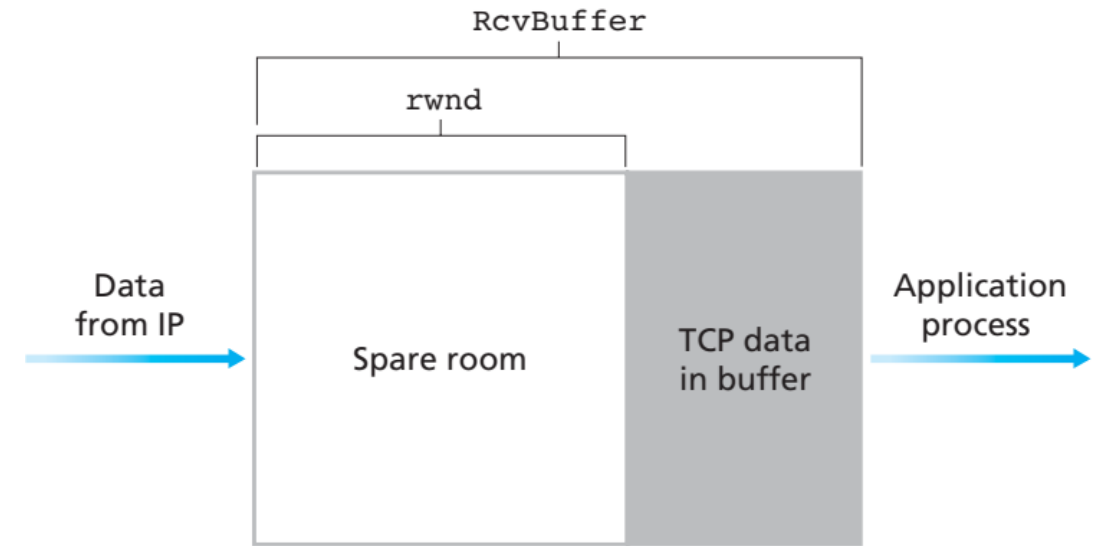


Receive Window (*rwnd*) and Receive Buffer (*RcvBuffer*)

- Flow control variables at the receiver end:
 - *RcvBuffer*: size of the **buffer** at the receiving host
 - *rwnd*: amount of **spare space** in the **receive buffer**; **dynamic** as the **spare space** in the **buffer changes** with **time**
 - *LastByteRead*: the **number of the last byte** in the **data stream** read from the **receive buffer** by the **receiver application process**
 - *LastByteRcvd*: the **number of the last byte** in the **data stream** that has been placed in the **receive buffer** at the receiver
- To prevent overflow of the allocated buffer:
 - $LastByteRcvd - LastByteRead \leq RcvBuffer$
 - $rwnd = RcvBuffer - [LastByteRcvd - LastByteRead]$

TCP: Flow Control

- Flow control variables at the sender end:
 - *LastByteSent*
 - *LastByteACKed*
- Amount of unacknowledged data that sender has sent into the connection:
 - $\text{LastByteSent} - \text{LastByteACKed}$
- TCP keeps the amount of unacknowledged data less than the value of *rwnd*
 - $\text{LastByteSent} - \text{LastByteACKed} \leq \text{rwnd}$
 - This ensures that the sender does not overflow the receiver buffer
- Technical challenge with receive window
 - Receiver's receive buffer becomes full ($\text{rwnd} = 0$)
 - After advertising $\text{rwnd} = 0$ to the sender, suppose the receiver has no data to sent (so no ACK segment)
 - TCP does not send new *rwnd* value to the sender even if the application process empties the receive buffer
 - Sender is blocked and cannot send more data



Receive Window (*rwnd*) and Receive Buffer (*RcvBuffer*)

- Solution:
 - Sender continues to send segments with one data byte when receiver's receive window is zero
 - Receiver acknowledges these segments and inserts the current value of the receive window
 - Eventually the buffer will begin to empty and the ACKs will contain a non-zero *rwnd* value