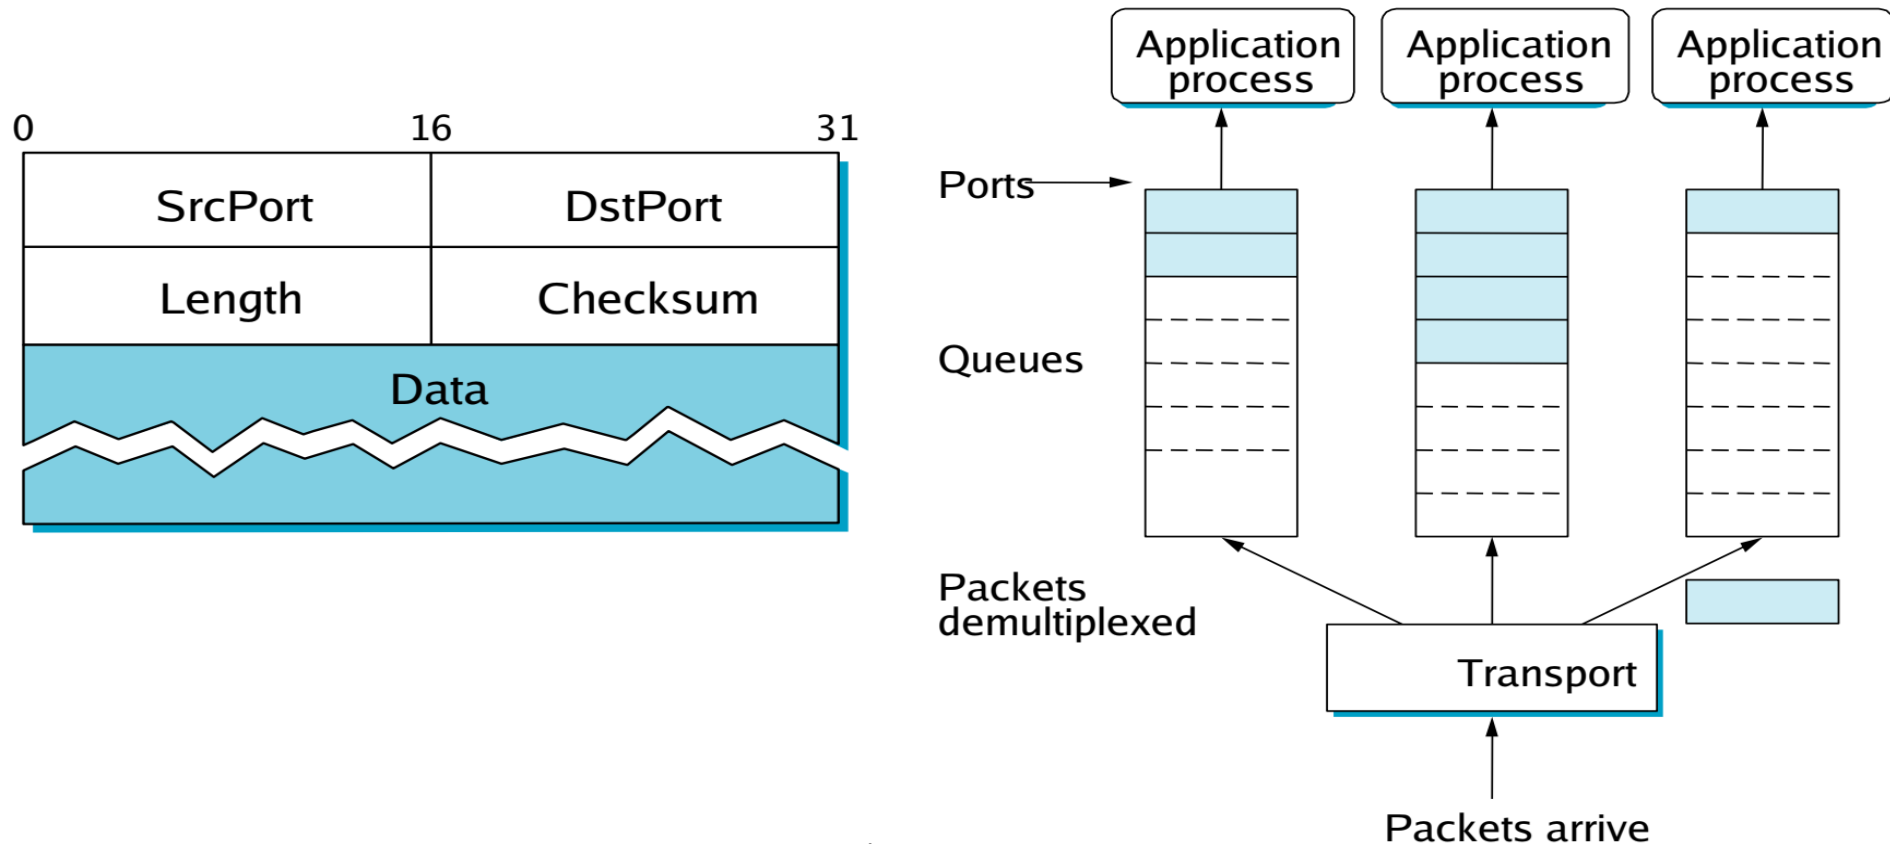


# TCP/IP Sockets

# Demultiplexing

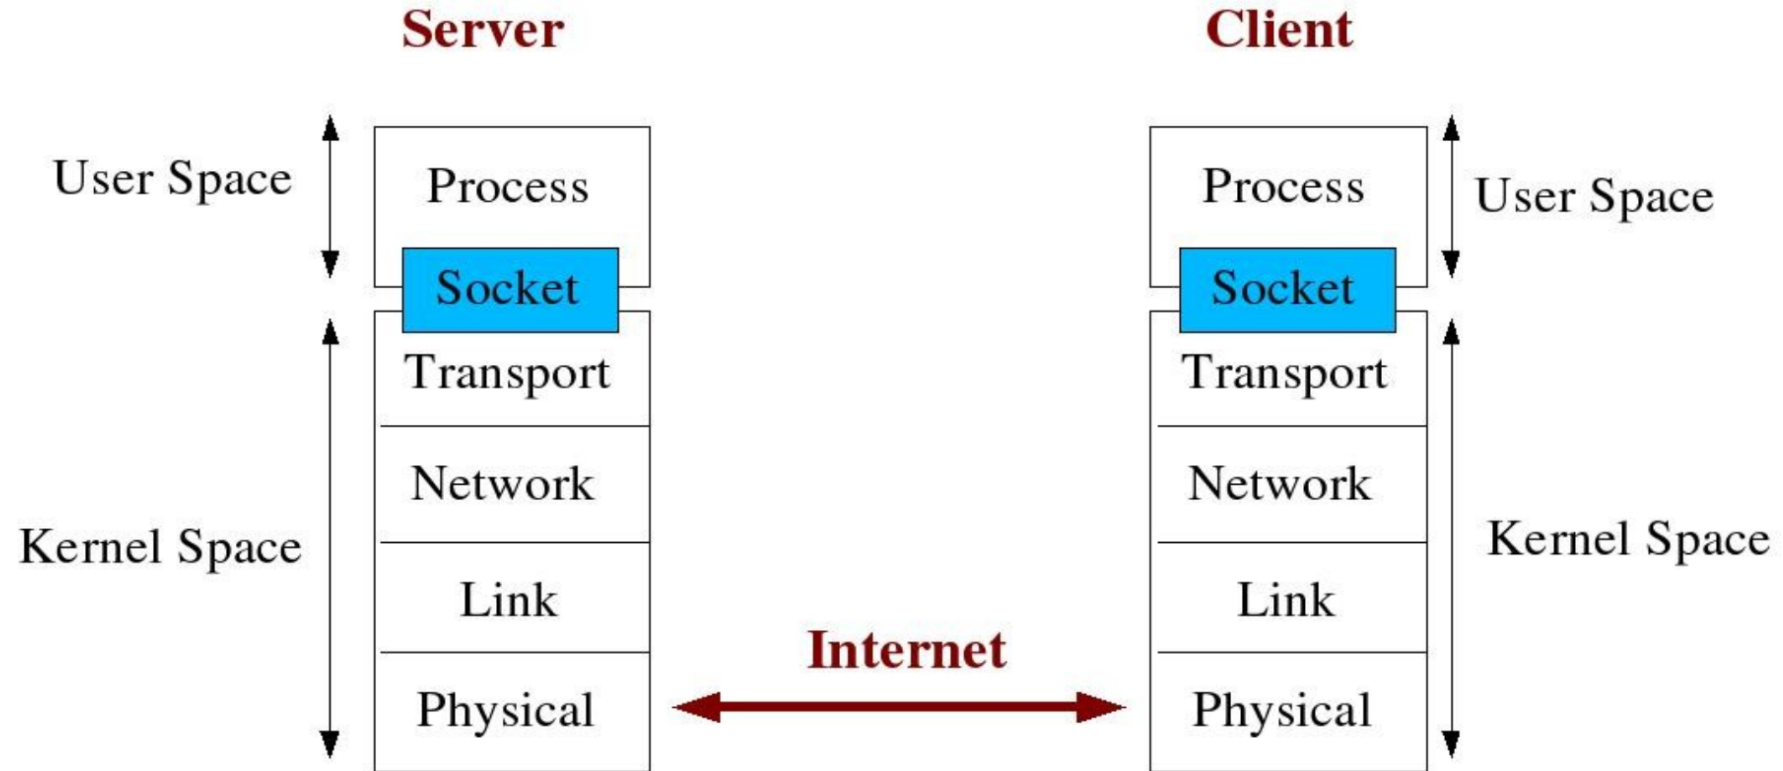
- Convert host-to-host packet delivery service into a process-to-process communication channel



# What is a socket?

- Socket: An interface between an application process and transport layer
  - The application process can send/receive messages to/from another application process (local or remote) via a socket
- In Unix jargon, a socket is a file descriptor – an integer associated with an open file
- Types of Sockets: **Internet Sockets**, unix sockets, X.25 sockets etc
  - Internet sockets characterized by IP Address (4 bytes), port number (2 bytes)

# Socket Description



# Types of Internet Sockets

- Stream Sockets (SOCK\_STREAM)
  - Connection oriented
  - Rely on TCP to provide reliable two-way connected communication
- Datagram Sockets (SOCK\_DGRAM)
  - Rely on UDP
  - Connection is unreliable

# Byte Ordering

- Two types of “Byte ordering”
  - Network Byte Order: High-order byte of the number is stored in memory at the lowest address
  - Host Byte Order: Low-order byte of the number is stored in memory at the lowest address
  - Network stack (TCP/IP) expects Network Byte Order
- Conversions:
  - htons() - Host to Network Short
  - htonl() - Host to Network Long
  - ntohs() - Network to Host Short
  - ntohl() - Network to Host Long

# socket() -- Get the file descriptor

- `int socket(int domain, int type, int protocol);`
  - domain should be set to `PF_INET`
  - type can be `SOCK_STREAM` or `SOCK_DGRAM`
  - set protocol to 0 to have socket choose the correct protocol based on type
  - `socket()` returns a socket descriptor for use in later system calls or -1 on error

```
int sockfd;
```

```
sockfd = socket (PF_INET, SOCK_STREAM, 0);
```

# Socket Structures

- struct sockaddr: Holds socket address information for many types of sockets

```
struct sockaddr {  
    unsigned short  sa_family;    //address family AF_XXX  
    unsigned short  sa_data[14]; //14 bytes of protocol addr  
}
```

- struct sockaddr\_in: A parallel structure that makes it easy to reference elements of the socket address

```
struct sockaddr_in {  
    short int          sin_family;    // set to AF_INET  
    unsigned short int sin_port;      // Port number  
    struct in_addr     sin_addr;      // Internet address  
    unsigned char       sin_zero[8];  //set to all zeros  
}
```

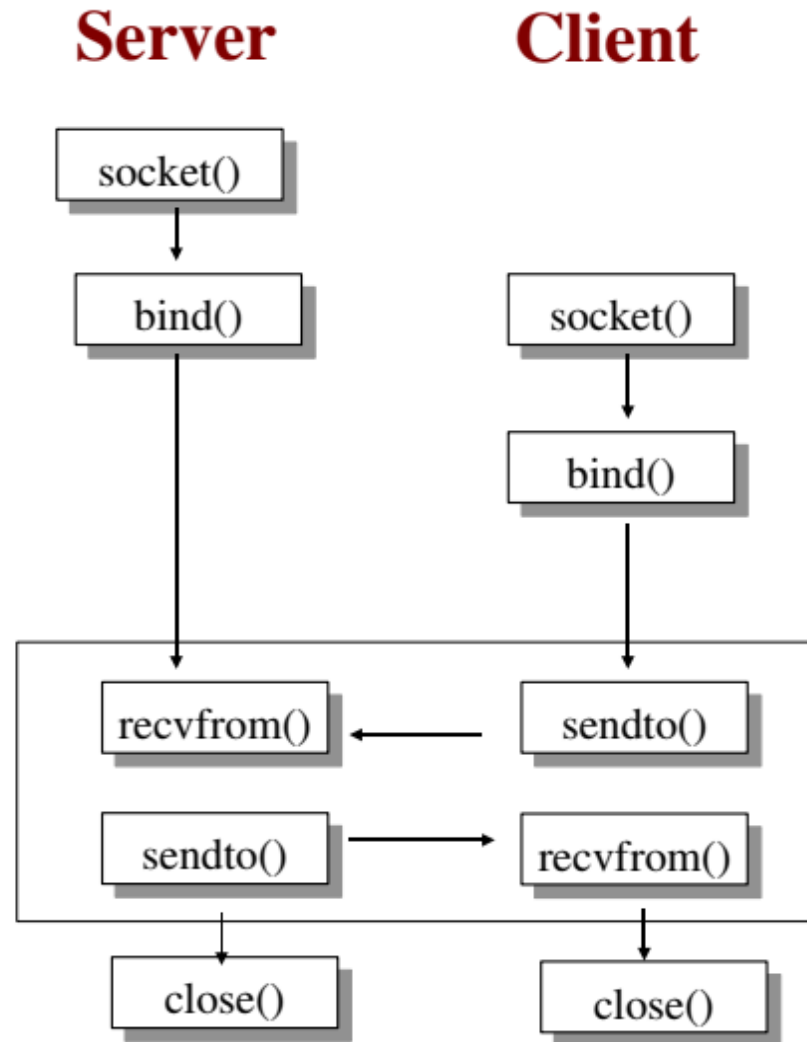
- sin\_port and sin\_addr must be in **Network Byte Order**



# bind() - what port am I on?

- Used to associate a socket with a port on the local machine
  - The port number is used by the kernel to match an incoming packet to a process
- `int bind(int sockfd, struct sockaddr *my_addr, int addrlen)`
  - `sockfd` is the socket descriptor returned by `socket()`
  - `my_addr` is pointer to `struct sockaddr` that contains information about your IP address and port
  - `addrlen` is set to `sizeof(struct sockaddr)`
  - returns -1 on error
  - `my_addr.sin_port = 0; //choose an unused port at random`
  - `my_addr.sin_addr.s_addr = INADDR_ANY; //use my IP adr`

# Connectionless Protocol



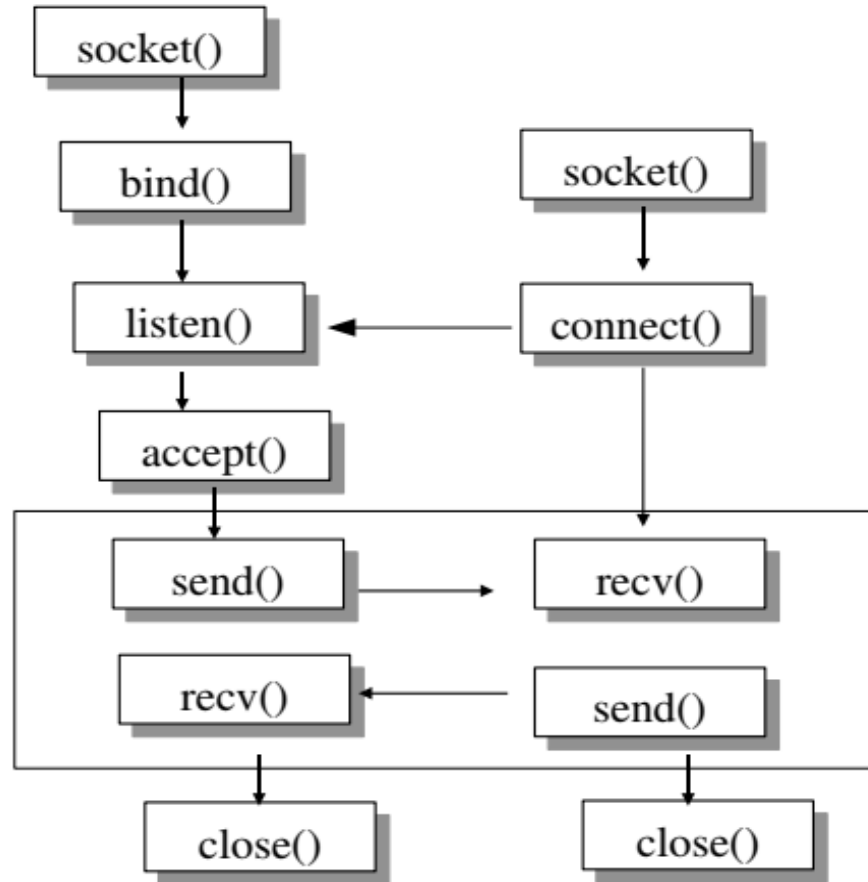
# sendto() and recvfrom() - DGRAM style

- `int sendto(int sockfd, const void *msg, int len, int flags, const struct sockaddr *to, int tolen);`
  - *to* is a pointer to a struct `sockaddr` which contains the destination IP and port
  - *tolen* is `sizeof(struct sockaddr)`
- `int recvfrom(int sockfd, void *buf, int len, int flags, struct sockaddr *from, int *fromlen);`
  - *from* is a pointer to a local struct `sockaddr` that will be filled with IP address and port of the originating machine
  - *fromlen* will contain length of address stored in *from*

# Connection Oriented Protocol

**Server**

**Client**



# connect() - Hello!

- Connects to a remote host
- `int connect(int sockfd, struct sockaddr *serv_addr, int addrlen)`
  - `sockfd` is the socket descriptor returned by `socket()`
  - `serv_addr` is pointer to `struct sockaddr` that contains information on destination IP address and port
  - `addrlen` is set to `sizeof(struct sockaddr)`
  - returns -1 on error
- No need to `bind()`, kernel will choose a port

# listen() - Call me please!

- Waits for incoming connections
- `int listen(int sockfd, int backlog);`
  - `sockfd` is the socket file descriptor returned by `socket()`
  - `backlog` is the number of connections allowed on the incoming queue
  - `listen()` returns -1 on error
  - Need to call `bind()` before you can `listen()`
    - `socket()`
    - `bind()`
    - `listen()`
    - `accept()`

# accept() - Thank you for calling !

- accept() gets the pending connection on the port you are listen()ing on
- `int accept(int sockfd, void *addr, int *addrlen);`
  - sockfd is the listening socket descriptor
  - information about incoming connection is stored in addr which is a pointer to a local struct `sockaddr_in`
  - addrlen is set to `sizeof(struct sockaddr_in)`
  - accept returns *a new socket file descriptor* to use for this accepted connection and -1 on error

# send() and recv() - Let's talk!

- The two functions are for communicating over stream sockets or connected datagram sockets.
- `int send(int sockfd, const void *msg, int len, int flags);`
  - `sockfd` is the socket descriptor you want to send data to (returned by `socket()` or got from `accept()`)
  - `msg` is a pointer to the data you want to send
  - `len` is the length of that data in bytes
  - set `flags` to 0 for now
  - `send()` returns the number of bytes actually sent (may be less than the number you told it to send) or -1 on error



# send() and recv() - Let's talk!

- `int recv(int sockfd, void *buf, int len, int flags);`
  - `sockfd` is the socket descriptor to read from
  - `buf` is the buffer to read the information into
  - `len` is the maximum length of the buffer
  - set flags to 0 for now
  - `recv()` returns the number of bytes actually read into the buffer or -1 on error
  - If `recv()` returns 0, the remote side has closed connection on you

# close() - Bye Bye!

- `int close(int sockfd);`
  - Closes connection corresponding to the socket descriptor and frees the socket descriptor
  - Will prevent any more sends and recvs