

TCP/IP Sockets - 1

Monday, February 14, 2022 3:26 PM

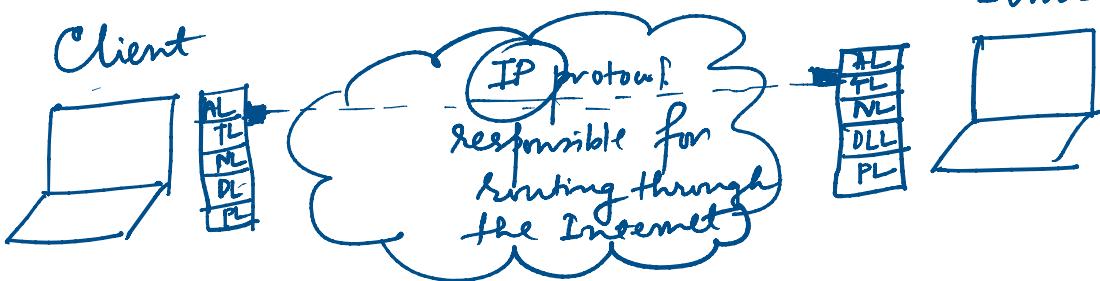
Open file — it can be a network connection, a FIFO, a pipe, a terminal and so on.

Socket — an integer associated with an open file.

UDP segment header — Src port, Dst port

TCP segment header — Src IP address, Src port,
Dst IP address, Dst port.

① Stream sockets — uses TCP to ensure that the data arrives sequentially and error-free and IP for Internet routing.



② Datagram sockets: uses UDP for data delivery and IP for Internet routing.

— Unreliable: the datagram which is sent may arrive out-of-order

— Connectionless: don't have to maintain an open connection as with stream sockets — just build a segment, ...

— Connectionless: done ...
 as with stream sockets — just build a segment,
 put an IP header (datagram/packet) on it with
 the destination information and send it out —
 no connection is needed.

Byte order:

The order in which bytes are stored in the memory —
 can be different and is an attribute of the processor.

Two common ordering systems:

① Little-Endian: low-order byte is stored at a lower address. This is also called Intel order since Intel's x86 family of CPUs popularized this ordering. Intel, AMD, VIA etc. are little-endian systems.

0	1	2	3	4	5	6	7	Address	Little Endian System.
Byte0	Byte1	Byte2	Byte3	Byte4	Byte5	Byte6	Byte7		

② Big-Endian: high-order byte is stored at a lower address. This is also called Motorola order since Motorola's Power PC architecture used this ordering. Motorola 68K and IBM Mainframes are big-endian systems.

Address	0	1	2	3	4	5	6	7
Bytes	Byte7	Byte6	Byte5	Byte4	Byte3	Byte2	Byte1	Byte0

Memory Content:

0x11	0x22	0x33	0x44	0x55	0x66	0x77	0x88
------	------	------	------	------	------	------	------

64-bit value on Little Endian:

0x8877665544332211

least significant byte
("octet") first

↳ Host Byte Ordering
(locally)

64-bit value on Big Endian:

0x1122334455667788

most significant byte ("octet")
first

→ Network Byte
Ordering
(Over the network)

To allow machines with different byte order conventions communicate with each other, the Internet protocols specify a canonical byte order convention in big-endian format for data transmitted over the network. This is known as Network Byte Order.

The ordering used on a computer is called Host Byte Order.

* A host system may be little-endian but while sending data into the n/w, it must convert data into big-endian format. Likewise, a little-endian machine must first convert network data into little-endian format before processing it.

Struct Sockaddr

→ feeds data into a socket

No init

Struct sockaddr

- sa-family - Value is:

AF_INET

↳ it is an address family that is used to designate the type of addresses that your socket can communicate with (e.g. IPv4 addresses)

- sa-data: contains the destination IP address and port number for the socket.

Struct in-addr {

unsigned long s-addr; // 32-bit long or 4 bytes in Network Byte Order.
};

sin-addr

and

sin-port

→ Network Byte Order - get encapsulated in the packet at the IP and Transport layers, respectively.

sin-family

→ Host Byte Order

↳ only used by the local kernel to determine what type of address the structure (sockaddr-in) contains and is not sent out on the network.

Sendto()

Returns -1 if there is an error.

int sendto (int sockfd, const void *msg, int len, unsigned int flags, const struct sockaddr *to, int tolen)

- `int sendto (int sockfd, const void *msg, int len, int flags, const struct sockaddr *to, int tolen)`
 - `flags`, `const struct sockaddr *to`, `int tolen`)
 - sends the no. of bytes actually sent
- `sockfd`: socket descriptor you want to send data to.
- `msg`: pointer to the data you want to send.
- `len`: length of the data in bytes.
- `flags`: set to 0.
- `to`: is a pointer to "struct sockaddr" that contains a destination IP address and port.
- `tolen`: simply set to `sizeof (struct sockaddr)`

recvfrom()

- `int recvfrom (int sockfd, void *buff, int len, unsigned int flags, struct sockaddr *from, int *fromlen)`
 - `>Returns -1 if there is an error.`
 - `int flags, struct sockaddr *from, int *fromlen`)
 - returns the no. of bytes received.
- `sockfd` - socket description to read data from.
- `buff` - buffer is to read the information into.
- `len` - maximum length of the buffer.
- `flags` - can be set to 0.
- `from` - pointer to the local "struct sockaddr" that will be filled with the IP address and port of the originating machine (source)
- `fromlen`: pointer to a local 'int' that should be initialized to `sizeof (struct sockaddr)`.

`connect()` - Connect to a remote host

connect() — Connect to a remote host

int connect (int sockfd, struct sockaddr *serv-addr, int addrlen);

Client side
returns -1 if there is an error.

Socket descriptor returned by socket()

pointer to "struct sockaddr" containing destination IP address and port no. set to sizeof(struct sockaddr).

- No need to call bind() — We don't care our local port numbers — the kernel which choose a local port number on behalf of us and the site we connect will automatically get this information from us.